

# Exploratory Data Analysis(EDA) & Regression using Ames Housing Data

## Objectives:

- To analyze and investigate the Ames Housing dataset and summarize its main characteristics.
- To build machine learning models to predict prices of houses

## Imports and Reading Data

### Importing Libraries:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns

%matplotlib inline
plt.style.use('ggplot') #default style for all visualisations
pd.set_option('display.max_columns', 100) #expanding no. of columns shown
pd.set_option('display.max_rows', 100) #expanding no. of rows shown
```

### About the dataset:

The Ames Housing dataset contains information about individual residential property in Ames, Iowa, from 2006 to 2010. The dataset was collected by Dean De Cock in 2011, and additional information is available via the following links:

- [A report describing the dataset](#)
- [Detailed documentation regarding the dataset's features](#)
- [The dataset in a tab-separated format](#)

### Reading the data:

```
In [2]: # Loading the dataset into a dataframe
url = "https://jse.amstat.org/v19n3/decock/AmesHousing.txt"
df = pd.read_csv(url, delimiter='\t')
```

## Understanding the dataset

An Overview:

```
In [3]: df.shape
```

Out[3]: (2930, 82)

```
In [4]: df.head()
```

Out[4]:

	Order	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilit
0	1	526301100	20	RL	141.0	31770	Pave	NaN	IR1	Lvl	AllF
1	2	526350040	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllF
2	3	526351010	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllF
3	4	526353030	20	RL	93.0	11160	Pave	NaN	Reg	Lvl	AllF
4	5	527105010	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllF

```
In [5]: df.tail()
```

Out[5]:

	Order	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	L
2925	2926	923275080	80	RL	37.0	7937	Pave	NaN	IR1	Lvl	
2926	2927	923276100	20	RL	NaN	8885	Pave	NaN	IR1	Low	
2927	2928	923400125	85	RL	62.0	10441	Pave	NaN	Reg	Lvl	
2928	2929	924100070	20	RL	77.0	10010	Pave	NaN	Reg	Lvl	
2929	2930	924151050	60	RL	74.0	9627	Pave	NaN	Reg	Lvl	

```
In [6]: #listing all columns at a glance
df.columns
```

```
Out[6]: Index(['Order', 'PID', 'MS SubClass', 'MS Zoning', 'Lot Frontage', 'Lot Area',  
             'Street', 'Alley', 'Lot Shape', 'Land Contour', 'Utilities',  
             'Lot Config', 'Land Slope', 'Neighborhood', 'Condition 1',  
             'Condition 2', 'Bldg Type', 'House Style', 'Overall Qual',  
             'Overall Cond', 'Year Built', 'Year Remod/Add', 'Roof Style',  
             'Roof Matl', 'Exterior 1st', 'Exterior 2nd', 'Mas Vnr Type',  
             'Mas Vnr Area', 'Exter Qual', 'Exter Cond', 'Foundation', 'Bsmt Qual',  
             'Bsmt Cond', 'Bsmt Exposure', 'BsmtFin Type 1', 'BsmtFin SF 1',  
             'BsmtFin Type 2', 'BsmtFin SF 2', 'Bsmt Unf SF', 'Total Bsmt SF',  
             'Heating', 'Heating QC', 'Central Air', 'Electrical', '1st Flr SF',  
             '2nd Flr SF', 'Low Qual Fin SF', 'Gr Liv Area', 'Bsmt Full Bath',  
             'Bsmt Half Bath', 'Full Bath', 'Half Bath', 'Bedroom AbvGr',  
             'Kitchen AbvGr', 'Kitchen Qual', 'TotRms AbvGrd', 'Functional',  
             'Fireplaces', 'Fireplace Qu', 'Garage Type', 'Garage Yr Blt',  
             'Garage Finish', 'Garage Cars', 'Garage Area', 'Garage Qual',  
             'Garage Cond', 'Paved Drive', 'Wood Deck SF', 'Open Porch SF',  
             'Enclosed Porch', '3Ssn Porch', 'Screen Porch', 'Pool Area', 'Pool QC',  
             'Fence', 'Misc Feature', 'Misc Val', 'Mo Sold', 'Yr Sold', 'Sale Type',  
             'Sale Condition', 'SalePrice'],  
            dtype='object')
```

```
In [7]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2930 entries, 0 to 2929

Data columns (total 82 columns):

#	Column	Non-Null Count	Dtype
0	Order	2930 non-null	int64
1	PID	2930 non-null	int64
2	MS SubClass	2930 non-null	int64
3	MS Zoning	2930 non-null	object
4	Lot Frontage	2440 non-null	float64
5	Lot Area	2930 non-null	int64
6	Street	2930 non-null	object
7	Alley	198 non-null	object
8	Lot Shape	2930 non-null	object
9	Land Contour	2930 non-null	object
10	Utilities	2930 non-null	object
11	Lot Config	2930 non-null	object
12	Land Slope	2930 non-null	object
13	Neighborhood	2930 non-null	object
14	Condition 1	2930 non-null	object
15	Condition 2	2930 non-null	object
16	Bldg Type	2930 non-null	object
17	House Style	2930 non-null	object
18	Overall Qual	2930 non-null	int64
19	Overall Cond	2930 non-null	int64
20	Year Built	2930 non-null	int64
21	Year Remod/Add	2930 non-null	int64
22	Roof Style	2930 non-null	object
23	Roof Matl	2930 non-null	object
24	Exterior 1st	2930 non-null	object
25	Exterior 2nd	2930 non-null	object
26	Mas Vnr Type	2907 non-null	object
27	Mas Vnr Area	2907 non-null	float64
28	Exter Qual	2930 non-null	object
29	Exter Cond	2930 non-null	object
30	Foundation	2930 non-null	object
31	Bsmt Qual	2850 non-null	object
32	Bsmt Cond	2850 non-null	object
33	Bsmt Exposure	2847 non-null	object
34	BsmtFin Type 1	2850 non-null	object
35	BsmtFin SF 1	2929 non-null	float64
36	BsmtFin Type 2	2849 non-null	object
37	BsmtFin SF 2	2929 non-null	float64
38	Bsmt Unf SF	2929 non-null	float64
39	Total Bsmt SF	2929 non-null	float64
40	Heating	2930 non-null	object
41	Heating QC	2930 non-null	object
42	Central Air	2930 non-null	object
43	Electrical	2929 non-null	object
44	1st Flr SF	2930 non-null	int64
45	2nd Flr SF	2930 non-null	int64
46	Low Qual Fin SF	2930 non-null	int64
47	Gr Liv Area	2930 non-null	int64
48	Bsmt Full Bath	2928 non-null	float64
49	Bsmt Half Bath	2928 non-null	float64
50	Full Bath	2930 non-null	int64
51	Half Bath	2930 non-null	int64
52	Bedroom AbvGr	2930 non-null	int64
53	Kitchen AbvGr	2930 non-null	int64
54	Kitchen Qual	2930 non-null	object

```

55 TotRms AbvGrd      2930 non-null   int64
56 Functional         2930 non-null   object
57 Fireplaces         2930 non-null   int64
58 Fireplace Qu       1508 non-null   object
59 Garage Type        2773 non-null   object
60 Garage Yr Blt      2771 non-null   float64
61 Garage Finish      2771 non-null   object
62 Garage Cars        2929 non-null   float64
63 Garage Area        2929 non-null   float64
64 Garage Qual        2771 non-null   object
65 Garage Cond        2771 non-null   object
66 Paved Drive        2930 non-null   object
67 Wood Deck SF       2930 non-null   int64
68 Open Porch SF      2930 non-null   int64
69 Enclosed Porch     2930 non-null   int64
70 3Ssn Porch         2930 non-null   int64
71 Screen Porch       2930 non-null   int64
72 Pool Area          2930 non-null   int64
73 Pool QC            13 non-null     object
74 Fence              572 non-null    object
75 Misc Feature       106 non-null    object
76 Misc Val           2930 non-null   int64
77 Mo Sold            2930 non-null   int64
78 Yr Sold            2930 non-null   int64
79 Sale Type          2930 non-null   object
80 Sale Condition     2930 non-null   object
81 SalePrice          2930 non-null   int64
dtypes: float64(11), int64(28), object(43)
memory usage: 1.8+ MB

```

## General observations:

Looking at the dataframe general information:

- There are 82 columns/variables
- The dataset contains 2930 records/rows of data.
- Also, the type of data is heterogeneous: both numerical and categorical columns of data are available.
- Most of the columns are assigned the object datatype
- There are some null values in the dataset

## Data Preparation

Combining, cleansing, enriching and transforming the raw Ames housing data to make it usable for Exploratory Data Analysis(EDA) and regression.

## Dropping columns:

Focusing the analysis on a few house features of major interest because the dataset has 82 features which are too many to analyse at once.

```

In [8]: #dropping columns by commenting them out. The purpose is to keep track of those
# retained and those that have been dropped.
df=df[[

```

```
#'Order', 'PID', 'MS SubClass',
'MS Zoning', 'Lot Frontage', 'Lot Area',
#'Street', 'Alley', 'Lot Shape',
'Land Contour',
#'Utilities', 'Lot Config', 'Land Slope',
'Neighborhood',
#'Condition 1', 'Condition 2',
'Bldg Type',
#'House Style',
'Overall Qual', 'Overall Cond', 'Year Built',
#'Year Remod/Add', 'Roof Style',
#'Roof Matl', 'Exterior 1st', 'Exterior 2nd', 'Mas Vnr Type',
#'Mas Vnr Area', 'Exter Qual', 'Exter Cond', 'Foundation', 'Bsmt Qual',
#'Bsmt Cond', 'Bsmt Exposure', 'BsmtFin Type 1', 'BsmtFin SF 1',
#'BsmtFin Type 2', 'BsmtFin SF 2', 'Bsmt Unf SF',
'Total Bsmt SF',
#'Heating', 'Heating QC', 'Central Air', 'Electrical', '1st Flr SF',
#'2nd Flr SF', 'Low Qual Fin SF',
'Gr Liv Area',
#'Bsmt Full Bath', 'Bsmt Half Bath', 'Full Bath', 'Half Bath', 'Kitchen AbvGr',
'Bedroom AbvGr', 'Full Bath',
'Kitchen Qual', 'TotRms AbvGrd',
#'Functional', 'Fireplaces', 'Fireplace Qu', 'Garage Type', 'Garage Yr Blt',
#'Garage Finish', 'Garage Cars', 'Garage Area', 'Garage Qual',
#'Garage Cond', 'Paved Drive', 'Wood Deck SF', 'Open Porch SF',
#'Enclosed Porch', '3Ssn Porch', 'Screen Porch', 'Pool Area', 'Pool QC',
#'Fence', 'Misc Feature', 'Misc Val',
'Mo Sold', 'Yr Sold',
'Sale Type', 'Sale Condition', 'SalePrice']].copy()
```

- **Columns have been reduced from 82 columns down to 20 columns**

In [9]: `df.shape`

Out[9]: (2930, 20)

## Renaming columns:

Removing spaces between column names and giving variables clear names with reference to the [data description guide](#) of the Ames Housing Dataset:

```
In [10]: df = df.rename(columns={'MS Zoning':'MS_Zoning', 'Lot Frontage':'LotFrontage_ft',
    'Land Contour':'Land_Contour',
    #'Neighborhood',
    'Bldg Type':'Building_Type', 'Overall Qual':'Overall_Quality', 'Overall Cond':'Overall_Condition',
    'Year Built':'Year_Built', 'Total Bsmt SF':'TotalBsmt_sqft', 'Gr Liv Area':'GrLivArea',
    'Bedroom AbvGr':'Bedroom_AbvGr', 'Full Bath':'Full_Bath',
    'Kitchen Qual':'Kitchen_Quality', 'TotRms AbvGrd':'TotRms_AbvGrd', 'Moi Sold':'MonthSold',
    'Sale Type':'Sale_Type', 'Sale Condition':'Sale_Condition', 'SalePrice':'SalePrice'})
```

## Assigning appropriate datatypes:

Overall\_Quality, Overall\_Condition and Month\_Sold are categorical columns whose values are to be changed from numbers to more descriptive representative text. Which operation is not possible if they remain allocated to the int64 datatype.

```
In [11]: #changing from int64 to object datatype
df[['Overall_Quality', 'Overall_Condition', 'Month_Sold']] =\
df[['Overall_Quality', 'Overall_Condition', 'Month_Sold']].astype(object)
```

## Renaming categorical feature values:

Replacing the unique values of the columns with more descriptive values in reference to the official [Ames Housing Data Description Documentation](#)

```
In [12]: # MS_Zoning Column
df['MS_Zoning'] = df['MS_Zoning'].replace({
    'RL': '(Res) Low Density', #(Res) stands for residential
    'RH': '(Res) High Density',
    'FV': '(Res) Floating Village',
    'RM': '(Res) Medium Density',
    'C (all)': 'Commercial',
    'I (all)': 'Industrial',
    'A (agr)': 'Agriculture'})
```

```
In [13]: # Overall_Quality Column
df['Overall_Quality'] = df['Overall_Quality'].replace({
    10: 'Very Excellent',
    9: 'Excellent',
    8: 'Very Good',
    7: 'Good',
    6: 'Above Average',
    5: 'Average',
    4: 'Below Average',
    3: 'Fair',
    2: 'Poor',
    1: 'Very Poor'})
```

```
In [14]: #overall condition columns
df['Overall_Condition'] = df['Overall_Condition'].replace({
    10: 'Very Excellent',
    9: 'Excellent',
    8: 'Very Good',
    7: 'Good',
    6: 'Above Average',
    5: 'Average',
    4: 'Below Average',
    3: 'Fair',
    2: 'Poor',
    1: 'Very Poor'})
```

```
In [15]: #Building_Type Column
df['Building_Type'] = df['Building_Type'].replace({
    '1Fam': 'Single-family Detached',
    '2fmCon': 'Two-family Conversion',
    'Duplx': 'Duplex',
    'TwnhsE': 'TwnHs End Unit',
    'Twnhs': 'TwnHs Inside Unit'})
```

```
In [16]: #Sale_Type Column
df['Sale_Type'] = df['Sale_Type'].replace({
    'WD ': 'Conventional WD', # WD stands for warranty deed
```

```

'CWD': 'Cash WD',
'VWD': 'VA Loan WD',
'New': 'New on mkt',
'COD': 'Court Officer Deed',
'Con': 'Contract regular',
'ConLw': '(Con) LowDown payt&I', #Con stands for contract
'ConLI': '(Con) Low I', #I stands for Interest
'ConLD': 'Contract LowDown',
'Oth': 'Other'})

```

```

In [17]: #Transforming Kitchen Column
df['Kitchen_Quality'] = df['Kitchen_Quality'].replace({
    'Ex': 'Excellent',
    'Gd': 'Good',
    'TA': 'Typical',
    'Fa': 'Fair',
    'Po': 'Poor'})

```

```

In [18]: #Transforming Land Contour column
df['Land_Contour'] = df['Land_Contour'].replace({
    'Lvl': 'Near Flat/Level',
    'Bnk': 'Banked', #- Quick and significant rise from street grade to building
    'HLS': 'Hillside', #- Significant slope from side to side
    'Low': 'Depression'})

```

```

In [19]: #Transforming Neighborhood column
df['Neighborhood'] = df['Neighborhood'].replace({
    'Blmngtn': 'Bloomington Heights',
    'Blueste': 'Bluestem',
    'BrDale': 'Briardale',
    'BrkSide': 'Brookside',
    'ClearCr': 'Clear Creek',
    'CollgCr': 'College Creek',
    'Crawfor': 'Crawford',
    'Edwards': 'Edwards',
    'Gilbert': 'Gilbert',
    'Greens': 'Greens',
    'GrnHill': 'Green Hills',
    'IDOTRR': 'Iowa DOT and Rail Road',
    'Landmrk': 'Landmark',
    'MeadowV': 'Meadow Village',
    'Mitchel': 'Mitchell',
    'NAmes': 'North Ames',
    'NoRidge': 'Northridge',
    'NPKvill': 'Northpark Villa',
    'NridgHt': 'Northridge Heights',
    'NWAmes': 'Northwest Ames',
    'OldTown': 'Old Town',
    'SWISU': 'South & West of Iowa State University',
    'Sawyer': 'Sawyer',
    'SawyerW': 'Sawyer West',
    'Somerst': 'Somerset',
    'StoneBr': 'Stone Brook',
    'Timber': 'Timberland',
    'Veenker': 'Veenker'})

```

```

In [20]: #Transforming 'Month_Sold' column
df['Month_Sold'] = df['Month_Sold'].replace({
    1: 'Jan',

```



```
2: 'Feb',
3: 'Mar',
4: 'Apr',
5: 'May',
6: 'Jun',
7: 'Jul',
8: 'Aug',
9: 'Sep',
10: 'Oct',
11: 'Nov',
12: 'Dec'})
```

## Feature Engineering:

House age is a critical house feature to analyse as it highly influences the sale price. Though its not readily available in the dataset, it can be created using the readily available Year\_Sold & Year\_Built features.

```
In [21]: #finding how old the property was as of the sale date
df['Property_Age'] = df['Year_Sold'] - df['Year_Built']
```

```
In [22]: df['Property_Age']
```

```
Out[22]: 0      50
1      49
2      52
3      42
4      13
..
2925    22
2926    23
2927    14
2928    32
2929    13
Name: Property_Age, Length: 2930, dtype: int64
```

## Checking for null values:

```
In [23]: df.isna().sum()
```

```
Out[23]: MS_Zoning      0
LotFrontage_ft    490
LotArea_sqft      0
Land_Contour      0
Neighborhood      0
Building_Type     0
Overall_Quality    0
Overall_Condition  0
Year_Built        0
TotalBsmt_sqft    1
GroundLivArea_sqft 0
Bedroom_AbvGr     0
Full_Bath         0
Kitchen_Quality    0
TotRms_AbvGrd     0
Month_Sold        0
Year_Sold         0
Sale_Type         0
Sale_Condition     0
SalePrice_USD     0
Property_Age      0
dtype: int64
```

### OBSERVATION:

- The lotFrontage\_ft and TotalBsmt\_sqft are the only columns with null values.
- lotFrontage\_ft column has 490 missing values and GroundLivArea\_sqft column has 1 missing value

## Checking for duplicates:

```
In [24]: #count of duplicates
df.duplicated().sum()
```

```
Out[24]: 2
```

```
In [25]: #eliminating the duplicates by making the inverse of the duplicates the new work
df = df[~df.duplicated()].reset_index(drop=True).copy()
```

```
In [26]: df.shape
```

```
Out[26]: (2928, 21)
```

- 2 Duplicate rows have been removed thus reducing the row count from 2930 to 2928

### Data Overview after data transformation & data cleaning

```
In [27]: df.head()
```

Out[27]:

	MS_Zoning	LotFrontage_ft	LotArea_sqft	Land_Contour	Neighborhood	Building_Type	Ove
0	(Res) Low Density	141.0	31770	Near Flat/Level	North Ames	Single-family Detached	Ab
1	(Res) High Density	80.0	11622	Near Flat/Level	North Ames	Single-family Detached	
2	(Res) Low Density	81.0	14267	Near Flat/Level	North Ames	Single-family Detached	Ab
3	(Res) Low Density	93.0	11160	Near Flat/Level	North Ames	Single-family Detached	
4	(Res) Low Density	74.0	13830	Near Flat/Level	Gilbert	Single-family Detached	

## Exploratory Data Analysis(EDA)

Analyzing and investigating the Ames Housing dataset and summarizing its main characteristics.

### Objectives:

- To understand patterns within the Ames housing data
- To detect outliers or anomalous events
- To find interesting relations among the housing variables

## Summary Statistics


In [28]: `df.describe().T`

Out[28]:

	count	mean	std	min	25%	50%	75%
<b>LotFrontage_ft</b>	2438.0	69.230517	23.373933	21.0	58.00	68.0	80.
<b>LotArea_sqft</b>	2928.0	10148.768101	7882.487925	1300.0	7440.75	9436.5	11556.
<b>Year_Built</b>	2928.0	1971.348361	30.253979	1872.0	1954.00	1973.0	2001.
<b>TotalBsmt_sqft</b>	2927.0	1051.923129	440.328019	0.0	793.00	990.0	1302.
<b>GroundLivArea_sqft</b>	2928.0	1499.780738	505.650793	334.0	1126.00	1442.0	1743.
<b>Bedroom_AbvGr</b>	2928.0	2.853825	0.827739	0.0	2.00	3.0	3.
<b>Full_Bath</b>	2928.0	1.565915	0.552436	0.0	1.00	2.0	2.
<b>TotRms_AbvGrd</b>	2928.0	6.442964	1.573012	2.0	5.00	6.0	7.
<b>Year_Sold</b>	2928.0	2007.789617	1.316683	2006.0	2007.00	2008.0	2009.
<b>SalePrice_USD</b>	2928.0	180817.827186	79905.769953	12789.0	129500.00	160000.0	213500.
<b>Property_Age</b>	2928.0	36.441257	30.300296	-1.0	7.00	34.0	54.

In [29]: `fig = px.box(df, x='Year_Sold', y='SalePrice_USD', points='outliers', color='Year_Sold',`

```
fig.show()
title='Visual Summary of House Prices over the Years'
```



### OBSERVATIONS:

- All the years have scenarios where the houses being sold at extremely high prices. The biggest outlier is detected in 2007 as USD 755,000.
- There is not a lot of variation in the selling price of the different houses throughout the 5 years.
- Nearly all houses were sold for a price below USD 400,000

## Feature Understanding

Exploring and analysing each housing variable in the Ames housing dataset, separately.

```
In [30]: # Creating a grid of subplots
fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(10, 12),
                        gridspec_kw={'hspace': 0.4, 'wspace': 0.5})

# Plotting the Kernel Density Plots in each subplot
sns.kdeplot(df['LotFrontage_ft'], ax=axs[0, 0], fill=True, color='blue')
sns.kdeplot(df['LotArea_sqft'], ax=axs[0, 1], fill=True, color='green')
```

```

sns.kdeplot(df['Year_Built'], ax=axes[0, 2], fill=True, color='red')
sns.kdeplot(df['TotalBsmt_sqft'], ax=axes[1, 0], fill=True, color='purple')
sns.kdeplot(df['GroundLivArea_sqft'], ax=axes[1, 1], fill=True, color='green')
sns.kdeplot(df['Bedroom_AbvGr'], ax=axes[1, 2], fill=True, color='red')
sns.kdeplot(df['TotRms_AbvGrd'], ax=axes[2, 0], fill=True, color='blue')
sns.kdeplot(df['SalePrice_USD'], ax=axes[2, 1], fill=True, color='green')
sns.kdeplot(df['Property_Age'], ax=axes[2, 2], fill=True, color='red')

# Setting titles for each subplot
axes[0, 0].set_title('Lot Frontage (lft)')
axes[0, 1].set_title('Lot Area Size (sqft)')
axes[0, 2].set_title('Year Built')
axes[1, 0].set_title('Basement Area (sqft)')
axes[1, 1].set_title('Living Area (sqft)')
axes[1, 2].set_title('Bedrooms')
axes[2, 0].set_title('Total Rooms')
axes[2, 1].set_title('Sale Price (USD)')
axes[2, 2].set_title('Property Age (yrs)')

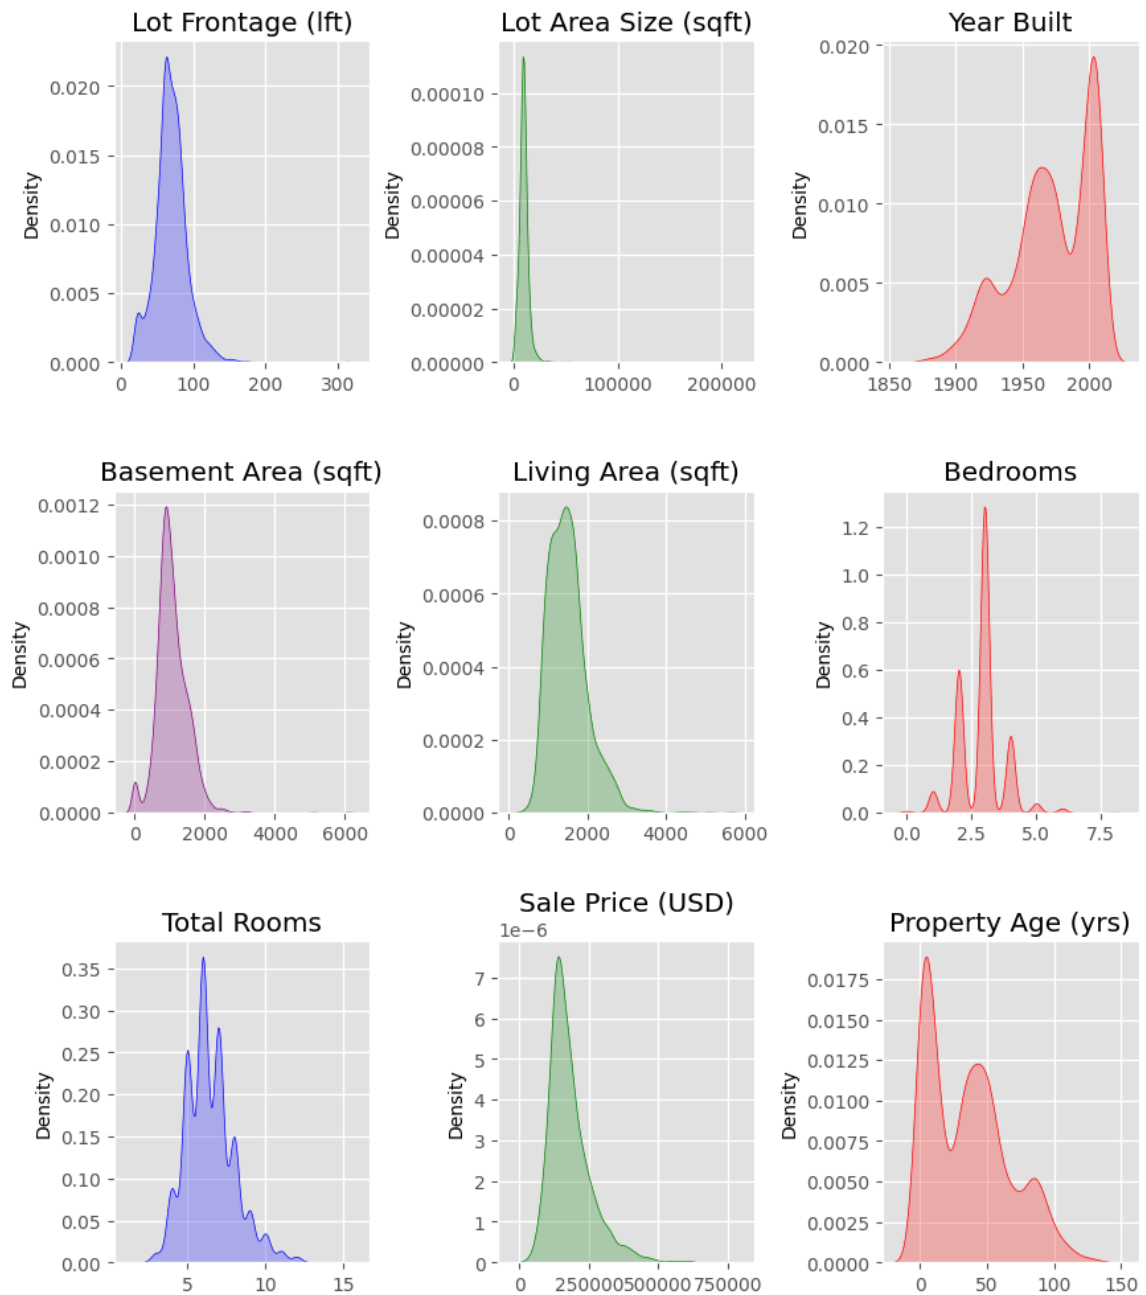
#Removing the x-axis label
axes[0, 0].set(xlabel='')
axes[0, 1].set(xlabel='')
axes[0, 2].set(xlabel='')
axes[1, 0].set(xlabel='')
axes[1, 1].set(xlabel='')
axes[1, 2].set(xlabel='')
axes[2, 0].set(xlabel='')
axes[2, 1].set(xlabel='')
axes[2, 2].set(xlabel='')

# Setting a title for the overall plot
fig.suptitle('DISTRIBUTION OF VARIOUS HOUSING VARIABLES', fontsize=16, fontweigh

# Displaying the plot
plt.show()

```

## DISTRIBUTION OF VARIOUS HOUSING VARIABLES



### OBSERVATIONS:

- Lot frontage, lot area size, basement area, living area and sale price features have a positively skewed distribution therefore most of the extreme values/outliers are on the right side thus they are higher values.
- Most of the houses have a basement size and living area size that ranges between 1000 - 2000 sqft

### Visualising categorical housing data

Neighbourhoods in the dataset are quite many, so it would make sense to visualize just a few of them, for example, the top 10. Therefore, below, a separate dataframe with top neighbourhoods from which the visualization is to be made has been created.

```
In [31]: # finding the top 10 neighbourhoods
df['Neighborhood'].value_counts().head(10)
```

```
Out[31]: North Ames          443
College Creek             267
Old Town                 239
Edwards                  193
Somerset                 182
Northridge Heights      166
Gilbert                  165
Sawyer                   151
Northwest Ames          131
Sawyer West              124
Name: Neighborhood, dtype: int64
```

```
In [32]: #creating dataframe with only top 10 neighbourhoods
top_10 = df.loc[df['Neighborhood'].isin(['College Creek', 'Old Town',
'Edwards', 'Somerset', 'Northridge Heights', 'Gilbert',
'Sawyer', 'Northwest Ames', 'Sawyer West'])]
```

```
In [33]: # Creating a grid of subplots
fig, axs = plt.subplots(nrows=4, ncols=3, figsize=(16, 14),
                        gridspec_kw={'hspace': 0.5, 'wspace': 0.5})

# Plotting the horizontal Count Plots in each subplot
ax1 = sns.countplot(data=df, y="MS_Zoning", ax=axs[0, 0],
                    order=df["MS_Zoning"].value_counts().index)
ax2 = sns.countplot(data=df, y="Land_Contour", ax=axs[0, 1])
ax3 = sns.countplot(data=top_10, y="Neighborhood", ax=axs[0, 2],
                    order=top_10["Neighborhood"].value_counts().index) #data source
ax4 = sns.countplot(data=df, y="Building_Type", ax=axs[1, 0],
                    order=df["Building_Type"].value_counts().index)
ax5 = sns.countplot(data=df, y="Overall_Quality", ax=axs[1, 1],
                    order=df["Overall_Quality"].value_counts().index)
ax6 = sns.countplot(data=df, y="Overall_Condition", ax=axs[1, 2],
                    order=df["Overall_Condition"].value_counts().index)
ax7 = sns.countplot(data=df, y="Kitchen_Quality", ax=axs[2, 0],
                    order=df["Kitchen_Quality"].value_counts().index)
ax8 = sns.countplot(data=df, y="Month_Sold", ax=axs[2, 1],
                    order=df["Month_Sold"].value_counts().index)
ax9 = sns.countplot(data=df, y="Sale_Type", ax=axs[2, 2],
                    order=df["Sale_Type"].value_counts().index)
ax10 = sns.countplot(data=df, y="Sale_Condition", ax=axs[3, 0],
                    order=df["Sale_Condition"].value_counts().index)
ax11 = sns.countplot(data=df, y="Full_Bath", ax=axs[3, 1],
                    order=df["Full_Bath"].value_counts().index)
ax12 = sns.countplot(data=df, y="Year_Sold", ax=axs[3, 2],
                    order=df["Year_Sold"].value_counts().index)

# Looping through each subplot and moving x-axis ticks to the top
for ax in [ax1, ax2, ax3, ax4, ax5, ax6,
           ax7, ax8, ax9, ax10, ax11, ax12]:
    ax.xaxis.set_ticks_position('top')

# Setting titles for each subplot
axs[0, 0].set_title('MS Zoning')
axs[0, 1].set_title('Land Contour')
axs[0, 2].set_title('Top 10 Neighbourhoods')
```

```

axs[1, 0].set_title('Building Type')
axs[1, 1].set_title('Overall Quality')
axs[1, 2].set_title('Overall Condition')
axs[2, 0].set_title('Kitchen Quality')
axs[2, 1].set_title('Month Sold')
axs[2, 2].set_title('Sale Type')
axs[3, 0].set_title('Sale Condition')
axs[3, 1].set_title('Full Bathrooms')
axs[3, 2].set_title('Year Sold')

#Looping through each subplot and removing the y-axis label
for ax in [ax1, ax2, ax2, ax3, ax4, ax5, ax6,
           ax7, ax8, ax9, ax10, ax11, ax12]:
    ax.set(ylabel='')

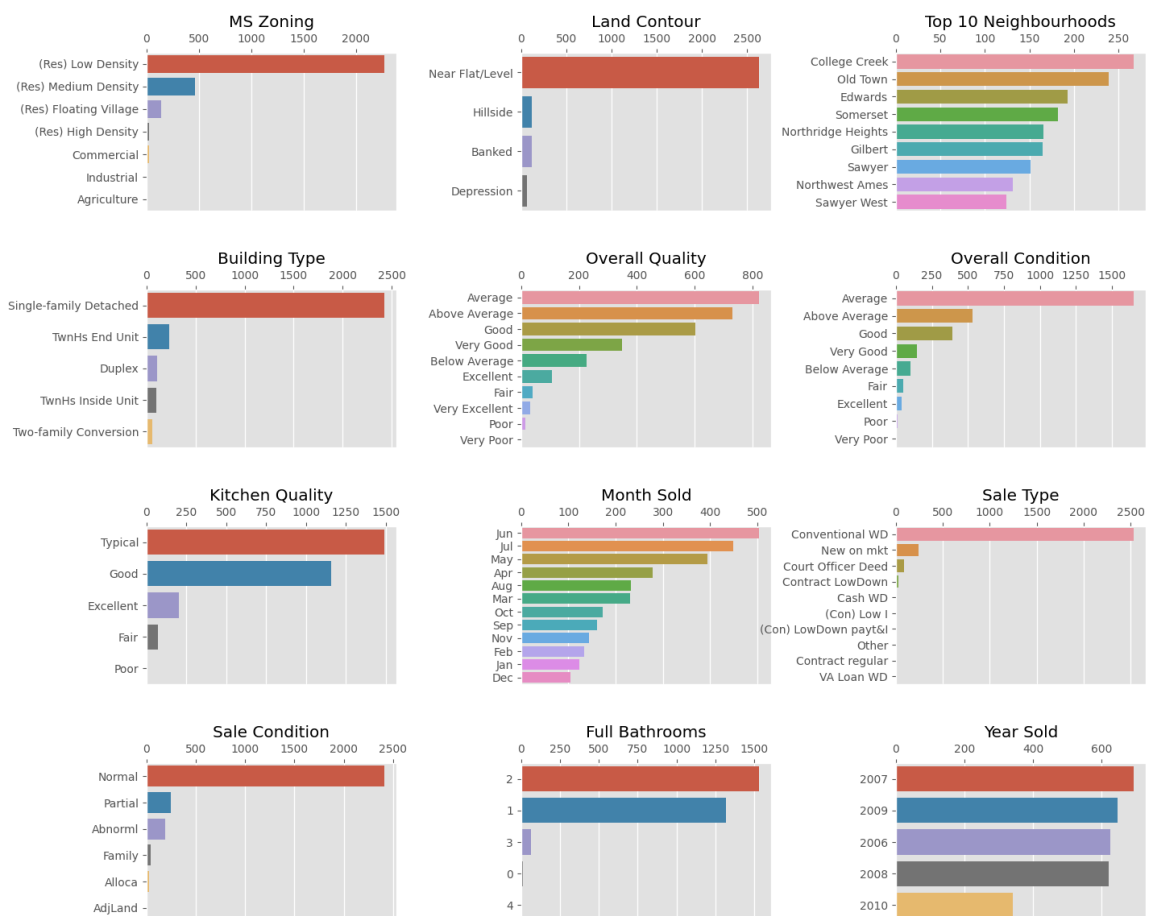
#Looping through each subplot and removing the x-axis label
for ax in [ax1, ax2, ax2, ax3, ax4, ax5, ax6,
           ax7, ax8, ax9, ax10, ax11, ax12]:
    ax.set(xlabel='')

# Setting a title for the overall plot
fig.suptitle('COUNTS OF OBSERVATIONS IN EACH CATEGORICAL BIN OF THE HOUSING VARI

# Displaying the plot
plt.show()

```

**COUNTS OF OBSERVATIONS IN EACH CATEGORICAL BIN OF THE HOUSING VARIABLES**



## OBSERVATIONS:

- Most of the houses sold were of average condition and quality



- More than 600 houses were sold in 2007, 2009, 2006 and also 2008
- Fewer houses were sold between January and December. This might probably be because of the holidays.
- Most individuals preferred to purchase houses in low density areas
- There was a low demand for Two-family conversion type of houses
- People in Ames mostly preferred purchasing houses neighbouring the Colege Creek area

## Feature Relationships

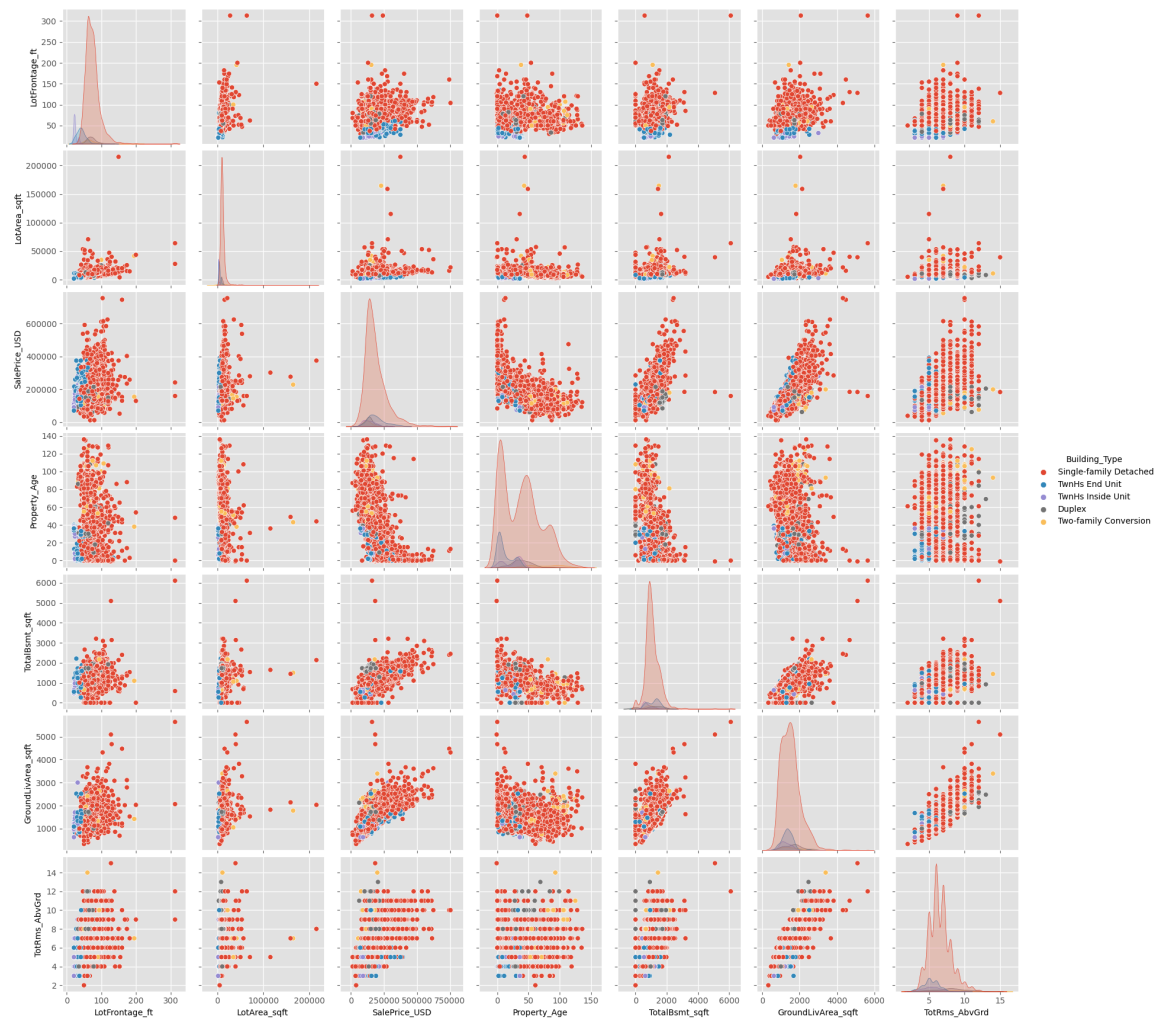
Analysing two quantitative housing variables to determine the nature of relationships between them.

```
In [34]: #creating a pairplot
g = sns.pairplot(df,
                 vars=['LotFrontage_ft', 'LotArea_sqft', 'SalePrice_USD', 'Property_
                    'TotalBsmt_sqft', 'GroundLivArea_sqft', 'TotRms_AbvGrd'],
                 hue='Building_Type')

# Adding overall title
g.fig.suptitle('A Visual Overview of the Various Feature Relationships', fontsize=14)

#showing the pairplot
plt.show()
```

## A Visual Overview of the Various Feature Relationships



### OBSERVATIONS:

- There is a positive relationship between:

Ground living Area and total number of rooms, Ground living Area and sale price ,  
Basement area and sale price

- There is a negative relationship between property age and sale price. As the property age of the house increases, the price for which its sold decreases
- There is no relationship between lot Frontage and the total rooms of a house. Therefore they have no measurable effect on each other

```
In [35]: #creating a pairplot
# Creating a grid of subplots
fig, axs = plt.subplots(nrows=2, ncols=3,
                        figsize=(8, 8), sharey=True,
                        gridspec_kw={'hspace': 0.2, 'wspace': 0.1})

ax1 = sns.kdeplot(
    data=df, x="LotFrontage_ft", y="SalePrice_USD",
    fill=True, thresh=0, levels=100, cmap="mako", ax=axs[0, 0])
ax2 = sns.kdeplot(
    data=df, x="Year_Sold", y="SalePrice_USD",
    fill=True, thresh=0, levels=100, cmap="mako", ax=axs[0, 1])
```

```

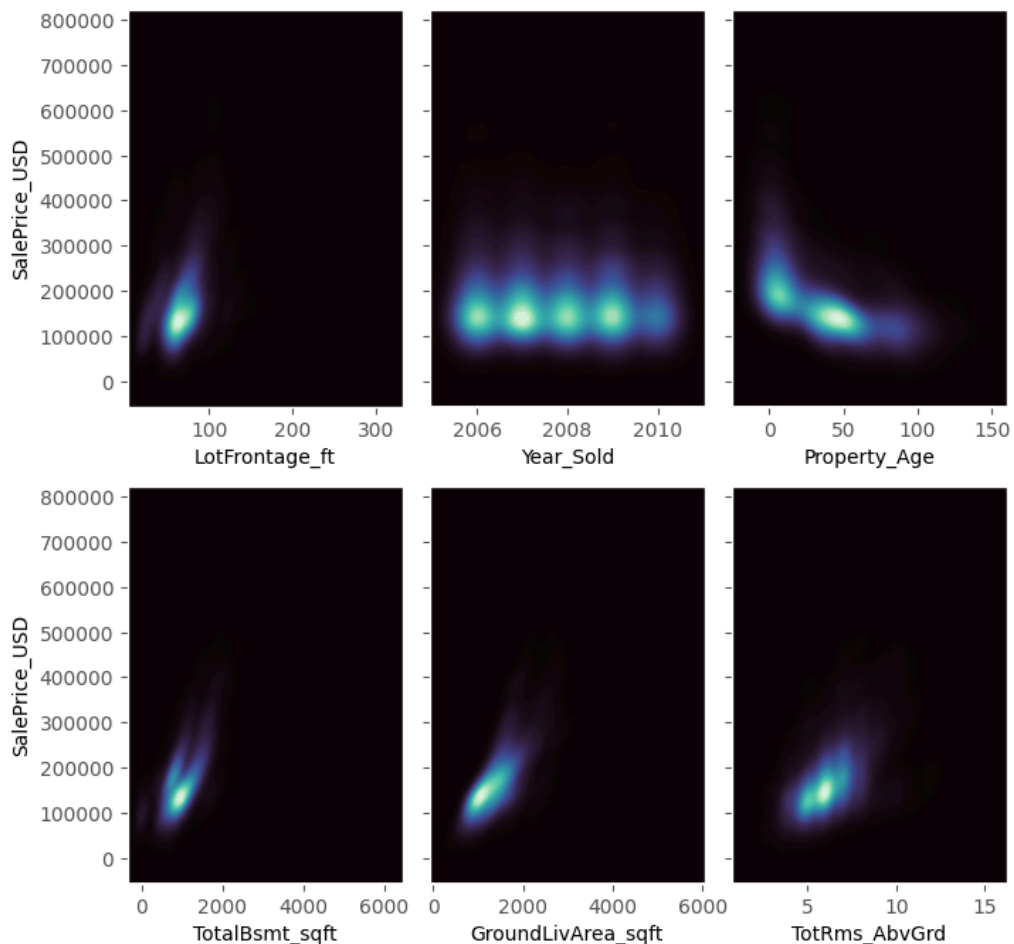
ax3 = sns.kdeplot(
    data=df, x="Property_Age", y="SalePrice_USD",
    fill=True, thresh=0, levels=100, cmap="mako", ax=axes[0, 2])
ax4 = sns.kdeplot(
    data=df, x="TotalBsm_t_sqft", y="SalePrice_USD",
    fill=True, thresh=0, levels=100, cmap="mako", ax=axes[1, 0])
ax5 = sns.kdeplot(
    data=df, x="GroundLivArea_sqft", y="SalePrice_USD",
    fill=True, thresh=0, levels=100, cmap="mako", ax=axes[1, 1])
ax6 = sns.kdeplot(
    data=df, x="TotRms_AbvGrd", y="SalePrice_USD",
    fill=True, thresh=0, levels=100, cmap="mako", ax=axes[1, 2])

# Adding overall title
fig.suptitle('A Visual Overview of Some Variables Affecting the Dwelling House P

#showing the pairplot
plt.show()

```

### A Visual Overview of Some Variables Affecting the Dwelling House Price



#### Observations:

- Lot frontage, total basement area, ground living are and total rooms of houses have a positive relationship with the sale price.
- The year a house is sold doesn't have a relationship with its price

- There's a negative relationship between the Sale price and property age

## CORRELATION:

Going a step further than just identifying if the relationship is positive or negative.


Below, the **extent** to which two housing variables are linearly related is measured:

```
In [36]: df_corr = df[['LotFrontage_ft', 'LotArea_sqft', 'Year_Built', 'TotalBsmt_sqft',
                    'GroundLivArea_sqft', 'Bedroom_AbvGr', 'Full_Bath', 'TotRms_AbvGrd',
                    'Year_Sold', 'SalePrice_USD', 'Property_Age']].dropna().corr()
df_corr
```

```
Out[36]:
```

	LotFrontage_ft	LotArea_sqft	Year_Built	TotalBsmt_sqft	GroundLivArea_sqft
LotFrontage_ft	1.000000	0.491849	0.122147	0.354031	0.38462
LotArea_sqft	0.491849	1.000000	0.061409	0.294912	0.34254
Year_Built	0.122147	0.061409	1.000000	0.435021	0.24278
TotalBsmt_sqft	0.354031	0.294912	0.435021	1.000000	0.46977
GroundLivArea_sqft	0.384629	0.342541	0.242788	0.469770	1.00000
Bedroom_AbvGr	0.241380	0.161172	-0.063066	0.052087	0.50731
Full_Bath	0.185950	0.149263	0.476621	0.346055	0.62405
TotRms_AbvGrd	0.354322	0.274950	0.116985	0.302828	0.81115
Year_Sold	-0.007405	-0.021224	-0.010509	-0.004913	-0.02951
SalePrice_USD	0.358163	0.321026	0.566885	0.647416	0.70609
Property_Age	-0.122296	-0.062217	-0.999122	-0.434653	-0.24370

```
In [37]: plt.figure(figsize=(8,6))
sns.heatmap(df_corr, annot=True, annot_kws={"fontsize":8},
            cmap="copper_r").set_title("Heatmap showing relationship between var")
plt.show()
```



#### Observations:

- There is a strong positive correlation(0.71) between the ground living area size of houses and their selling price. Therefore, the `GroundLivArea_sqft` feature can be considered one of the most important features for model training in a bid to predict house prices.

## REGRESSION

Building machine learning models to predict prices of houses based on their ground living area size.

### Linear Regression

- Splitting data into X and Y variables**

Separating input values (features) and the expected output (label) into separate numpy arrays

```
In [38]: y = df['SalePrice_USD']
y
```

```
Out[38]: 0      215000
         1      105000
         2      172000
         3      244000
         4      189900
         ...
        2923    142500
        2924    131000
        2925    132000
        2926    170000
        2927    188000
Name: SalePrice_USD, Length: 2928, dtype: int64
```

Below, a `reshape` on the input data is performed in order for the Linear Regression package to understand it correctly. Linear Regression expects a 2D-array as an input, where each row of the array corresponds to a vector of input features. In this case, since I have only one input - an array with shape  $N \times 1$  is needed, where  $N$  is the dataset size.

```
In [39]: X = df['GroundLivArea_sqft'].to_numpy().reshape(-1,1)
         X
```

```
Out[39]: array([[1656],
                [ 896],
                [1329],
                ...,
                [ 970],
                [1389],
                [2000]], dtype=int64)
```

- **Performing 80/20 data split**

Splitting the data into train and test datasets, so as to validate the model after training.

```
In [40]: from sklearn.model_selection import train_test_split

         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

Now the following four variables have been created:

- `X_train`: The `GroundLivArea_sqft` feature values to be used to train the model
- `y_train`: The corresponding `SalePrice_USD` labels to be used to train the model
- `X_test`: The `GroundLivArea_sqft` feature values to be used to validate the model
- `y_test`: The corresponding `SalePrice_USD` labels to be used to validate the model

- **Training the model**

Defining the `LinearRegression` object, and fitting it to the data using the `fit` method:

```
In [41]: from sklearn.linear_model import LinearRegression

         lin_reg = LinearRegression()
         lin_reg.fit(X_train,y_train)
```

Out[41]: ▾ LinearRegression  
LinearRegression()

- **Applying the model to make house sale price predictions**

```
In [42]: y_lr_train_pred = lin_reg.predict(X_train)

# predicting prices on a test dataset
y_lr_test_pred = lin_reg.predict(X_test)
```

```
In [43]: y_lr_train_pred
```

```
Out[43]: array([195202.24929397, 109967.13297948, 100811.44273598, ...,
                194112.28616975, 206537.86578593, 236402.85538973])
```

```
In [44]: y_lr_test_pred
```

```
Out[44]: array([314771.2040216 , 198581.13497908, 211551.69615737, 145063.94557957,
145935.91607895, 200543.06860268, 104408.32104592, 172531.01631007,
150513.7612007 , 111711.07397824, 227029.17252138, 109640.14404221,
233132.96601705, 157925.51044544, 199998.08704057, 146262.90501621,
150949.74645039, 111711.07397824, 179833.76924238, 182449.68074053,
154437.62844791, 231607.01764313, 201306.04278964, 123264.68309504,
193131.31935794, 148333.83495224, 241416.68576117, 195638.23454366,
172095.03106038, 193894.2935449 , 165991.23756471, 242179.65994813,
111711.07397824, 132747.36227581, 191169.38573434, 207409.83628531,
182558.67705295, 150949.74645039, 294824.87884826, 212205.6740319 ,
160977.40719327, 108332.18829314, 274987.54998734, 189316.44842315,
260055.05518544, 295696.84934764, 115416.94860061, 178307.82086847,
150186.77226343, 166100.23387713, 161413.39244296, 106806.23991922,
167408.1896262 , 246103.52719534, 147897.84970255, 181795.70286599,
228228.13195803, 241852.67101086, 86205.93687135, 208826.7883468 ,
164901.27444048, 162285.36294234, 105607.28048257, 210243.7404083 ,
213731.62240582, 165228.26337775, 131984.38808885, 178961.798743 ,
270627.69749043, 153565.65794853, 197055.18660516, 168934.13800012,
187572.50742439, 234767.91070339, 147243.87182802, 199889.09072815,
177871.83561878, 195856.22716851, 136344.24058575, 196619.20135547,
151821.71694977, 204575.93216232, 167953.17118832, 183757.6364896 ,
96996.57180119, 151385.73170008, 178525.81349331, 132965.35490065,
155527.59157214, 257112.15475003, 192259.34885856, 147897.84970255,
148333.83495224, 176236.89093244, 196946.19029274, 202178.01328902,
204793.92478717, 173402.98680945, 195202.24929397, 93617.68611608,
345072.17887509, 147788.85339013, 119558.80847267, 246103.52719534,
192368.34517099, 89475.82624402, 207191.84366046, 176236.89093244,
194112.28616975, 140050.11520812, 177653.84299393, 111711.07397824,
191169.38573434, 234658.91439097, 128932.49134102, 353028.90968194,
144191.97508018, 119558.80847267, 174928.93518337, 136562.2332106 ,
130131.45077767, 113673.00760185, 217655.48965304, 194766.26404428,
198254.14604181, 150513.7612007 , 278257.43936002, 254278.25062704,
238582.78163818, 164356.29287837, 322182.95326634, 206210.87684866,
288503.09272774, 191496.3746716 , 111711.07397824, 199780.09441572,
359786.68105214, 303108.59859238, 487966.34446115, 196728.19766789,
255586.20637611, 171659.04581069, 292208.96735011, 160977.40719327,
74216.34250486, 192041.35623372, 150513.7612007 , 126098.58721803,
112801.03710247, 143755.98983049, 130240.44709009, 90783.7819931 ,
179288.78768027, 172531.01631007, 163157.33344172, 275314.53892461,
136453.23689818, 298312.76084578, 305942.50271536, 179506.78030512,
387580.74071991, 197927.15710454, 145717.9234541 , 260927.02568482,
148333.83495224, 159887.44406904, 131330.41021431, 227901.14302076,
114980.96335092, 152148.70588704, 168607.14906285, 136562.2332106 ,
189316.44842315, 150840.75013796, 198472.13866665, 253079.29119039,
138851.15577148, 229754.08033195, 159451.45881935, 244250.58988416,
190188.41892253, 374937.16847889, 130894.42496462, 264087.91874508,
162503.35556719, 172531.01631007, 158688.48463239, 96015.60498938,
169261.12693739, 283816.25129357, 178743.80611816, 261363.01093451,
111711.07397824, 114217.98916396, 122174.71997081, 131439.40652674,
122610.7052205 , 221797.3495251 , 239236.75951272, 107460.21779376,
217764.48596546, 90783.7819931 , 169261.12693739, 169479.11956223,
206537.86578593, 203812.95797536, 195529.23823124, 253842.26537735,
163375.32606657, 214385.60028036, 310411.35152469, 193458.30829521,
122174.71997081, 166645.21543925, 177108.86143182, 167408.1896262 ,
99721.47961175, 186482.54430016, 178089.82824362, 131984.38808885,
181141.72499146, 132311.37702612, 230081.06926922, 152257.70219946,
245122.56038354, 150731.75382554, 236184.86276488, 164901.27444048,
200107.08335299, 274551.56473765, 199453.10547846, 171441.05318584,
126970.55771741, 165991.23756471, 176781.87249455, 132420.37333854,
183539.64386475, 136562.2332106 , 136562.2332106 , 150295.76857585,
```



353900.88018132, 198908.12391634, 172967.00155976, 272589.63111404,  
166863.20806409, 109095.1624801 , 136562.2332106 , 198036.15341696,  
114980.96335092, 368724.3786708 , 139723.12627086, 114435.98178881,  
76396.26875331, 189970.42629769, 121956.72734597, 167299.19331378,  
132311.37702612, 197818.16079212, 149532.79438889, 188662.47054862,  
242942.63413509, 196401.20873062, 236947.83695184, 214167.60765551,  
186700.53692501, 356734.78430431, 154982.61001002, 318150.0897067 ,  
188226.48529893, 236511.85170215, 146480.89764106, 174492.94993367,  
159669.4514442 , 172967.00155976, 204684.92847474, 219399.4306518 ,  
176236.89093244, 213513.62978098, 294170.90097372, 129368.47659071,  
153565.65794853, 175909.90199517, 213731.62240582, 305288.52484083,  
175800.90568275, 233568.95126674, 146153.90870379, 208826.7883468 ,  
226048.20570958, 292317.96366254, 146371.90132864, 163266.32975414,  
90129.80411856, 130894.42496462, 151603.72432492, 192913.3267331 ,  
181032.72867903, 136562.2332106 , 155527.59157214, 142012.04883173,  
195420.24191882, 195856.22716851, 186809.53323743, 105062.29892046,  
330575.66932288, 116942.89697453, 153783.65057338, 188226.48529893,  
160759.41456842, 134491.30327457, 174165.96099641, 173402.98680945,  
128714.49871617, 237056.83326426, 152693.68744915, 231171.03239344,  
188117.4889865 , 148333.83495224, 173729.97574672, 207409.83628531,  
146589.89395348, 141467.06726962, 116724.90434968, 147897.84970255,  
111711.07397824, 205992.88422381, 111711.07397824, 150295.76857585,  
181250.72130388, 194548.27141944, 155854.58050941, 130894.42496462,  
182885.66599022, 161849.37769265, 271717.66061466, 141794.05620689,  
205120.91372443, 175800.90568275, 208281.80678469, 118250.8527236 ,  
245776.53825807, 233459.95495432, 136562.2332106 , 226920.17620896,  
169915.10481192, 123700.66834473, 148660.82388951, 225503.22414746,  
200870.05753995, 227029.17252138, 153892.6468858 , 175800.90568275,  
196728.19766789, 176236.89093244, 276186.50942399, 152257.70219946,  
180160.75817965, 137107.21477271, 207191.84366046, 224849.24627293,  
206537.86578593, 196510.20504305, 231389.02501829, 325779.83157628,  
188117.4889865 , 172095.03106038, 125226.61671865, 204793.92478717,  
198036.15341696, 153783.65057338, 199562.10179088, 314880.20033402,  
267684.79705502, 149532.79438889, 203921.95428778, 250790.36862951,  
103645.34685897, 168171.16381316, 123046.69047019, 248828.43500591,  
125117.62040622, 130458.43971493, 218527.46015242, 218636.45646484,  
141249.07464477, 127406.5429671 , 189970.42629769, 111929.06660309,  
110839.10347886, 204357.93953747, 264196.9150575 , 196292.2124182 ,  
200652.06491511, 106479.25098195, 323817.89795268, 109640.14404221,  
179070.79505543, 254278.25062704, 217655.48965304, 247738.47188168,  
164247.29656595, 147897.84970255, 221143.37165056, 151494.7280125 ,  
180378.7508045 , 247847.4681941 , 139287.14102117, 111711.07397824,  
155854.58050941, 223323.29789901, 131875.39177643, 113346.01866458,  
130022.45446524, 412867.88520196, 134927.28852426, 194221.28248217,  
126098.58721803, 150077.77595101, 194657.26773186, 165555.25231502,  
150840.75013796, 307032.46583959, 120321.78265963, 201524.03541449,  
106479.25098195, 187354.51479954, 205011.91741201, 132638.36596339,  
287849.11485321, 164029.3039411 , 264305.91136992, 236511.85170215,  
120866.76422174, 93181.70086639, 114762.97072607, 159015.47356966,  
182994.66230264, 116942.89697453, 124790.63146896, 198472.13866665,  
65278.6448862 , 179942.76555481, 202178.01328902, 177980.8319312 ,  
223650.28683628, 140922.08570751, 134382.30696215, 181577.71024115,  
306378.48796505, 126425.5761553 , 156290.5657591 , 155745.58419698,  
229427.09139468, 130894.42496462, 116942.89697453, 122937.69415777,  
211551.69615737, 320003.02701788, 111711.07397824, 168389.15643801,  
230953.0397686 , 184956.59592625, 181250.72130388, 173838.97205914,  
97323.56073845, 216674.52284123, 201633.03172691, 111384.08504097,  
243814.60463447, 160541.42194358, 163266.32975414, 136562.2332106 ,  
248828.43500591, 173402.98680945, 111711.07397824, 201306.04278964,  
227247.16514623, 263869.92612023, 205883.88791139, 166645.21543925,

```
180487.74711692, 209262.77359649, 106479.25098195, 140268.10783297,
120103.79003478, 161304.39613054, 129913.45815282, 257330.14737487,
156399.56207152, 310738.34046196, 196510.20504305, 148333.83495224,
142448.03408142, 174056.96468398, 194984.25666913, 193022.32304552,
130458.43971493, 292426.95997496, 177544.84668151, 173947.96837156,
263433.94087054, 127079.55402983, 136562.2332106 , 202831.99116356,
179288.78768027, 90783.7819931 , 198145.14972939, 221252.36796298,
237819.80745122, 116942.89697453, 148333.83495224, 231171.03239344,
202396.00591387, 221034.37533814, 215039.57815489, 190297.41523496,
166536.21912682, 98958.50542479, 182122.69180326, 243705.60832204,
161849.37769265, 150186.77226343, 178743.80611816, 130894.42496462,
203049.9837884 , 259183.08468606, 245340.55300838, 119122.82322298,
199671.0981033 , 255150.22112642, 168934.13800012, 353900.88018132,
265395.87449415, 217982.4785903 , 156726.55100879, 125226.61671865,
181468.71392872, 125226.61671865, 173402.98680945, 219726.41958907,
197382.17554243, 172967.00155976, 114980.96335092, 135363.27377395,
201742.02803933, 179288.78768027, 154873.6136976 , 228337.12827045,
131330.41021431, 199562.10179088, 311610.31096134, 102991.36898443,
139614.12995843, 125117.62040622, 237601.81482638, 180160.75817965,
228882.10983257, 229863.07664437, 119558.80847267, 285887.1812296 ,
126970.55771741, 201742.02803933, 151058.74276281, 259837.06256059,
182231.68811568, 179833.76924238, 181141.72499146, 314662.20770917,
155854.58050941, 189098.45579831, 168171.16381316, 134382.30696215,
96015.60498938, 128714.49871617, 180378.7508045 , 184847.59961382,
212859.65190644, 111711.07397824, 243051.63044751, 186700.53692501,
172204.0273728 , 221906.34583752, 149859.78332616, 182449.68074053,
305288.52484083, 194766.26404428])
```

- **Evaluating model performance**

Measuring how close house price predictions are to the expected prices.

```
In [45]: from sklearn.metrics import mean_squared_error, r2_score
```

```
lr_train_mse = mean_squared_error(y_train, y_lr_train_pred)
lr_train_r2 = r2_score(y_train, y_lr_train_pred)

lr_test_mse = mean_squared_error(y_test, y_lr_test_pred)
lr_test_r2 = r2_score(y_test, y_lr_test_pred)
```

```
In [46]: print('LR MSE (Train): ', lr_train_mse)
print('LR R2 (Train): ', lr_train_r2)
print('LR MSE (Test): ', lr_test_mse)
print('LR R2 (Test): ', lr_test_r2)
```

```
LR MSE (Train): 3215818093.8122363
LR R2 (Train): 0.47869722670799963
LR MSE (Test): 3112173832.7581115
LR R2 (Test): 0.5696028228959853
```

```
In [47]: lr_results = pd.DataFrame(['Linear Regression', lr_train_mse, lr_train_r2,
                                   lr_test_mse, lr_test_r2]).transpose()

#rename columns
lr_results.columns = ['Method', 'Training MSE', 'Training R2', 'Test MSE', 'Test R2']

lr_results
```

Out[47]:


	Method	Training MSE	Training R2	Test MSE	Test R2
0	Linear Regression	3215818093.812236	0.478697	3112173832.758111	0.569603

Observation:

- The co-efficient of determination(Test R2) of the linear regression model is 0.57, a value that lies between 0 and 1. Therefore, this linear regression model partially predicts the price of houses. If the co-efficient of determination was 1, it could have meant that this model perfectly predicts the price of houses

### Plotting the test data together with the regression line to better evaluate the linear regression model performance

```
In [48]: plt.scatter(X_test,y_test,c="#4c2a4c", alpha=0.3)
plt.plot(X_test,y_lr_test_pred)
plt.show()
```



Observation: The regression line is closely following the scatter points. This suggests that the linear regression model is making accurate house price predictions.

## Random Forest Regression

- **Training the model**

```
In [49]: from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(max_depth=2, random_state=0)
rf.fit(X_train, y_train)
```

Out[49]:

```

RandomForestRegressor
RandomForestRegressor(max_depth=2, random_state=0)

```

- Applying the model to make house sale price predictions

In [50]:

```

y_rf_train_pred = rf.predict(X_train)
y_rf_test_pred = rf.predict(X_test)

```

- Evaluating model performance

In [51]:

```

from sklearn.metrics import mean_squared_error, r2_score

```

```

rf_train_mse = mean_squared_error(y_train, y_rf_train_pred)
rf_train_r2 = r2_score(y_train, y_rf_train_pred)

rf_test_mse = mean_squared_error(y_test, y_rf_test_pred)
rf_test_r2 = r2_score(y_test, y_rf_test_pred)

```

In [52]:

```

rf_results = pd.DataFrame(['Random Forest', rf_train_mse, rf_train_r2,
                           rf_test_mse, rf_test_r2]).transpose()

#rename columns
rf_results.columns = ['Method', 'Training MSE', 'Training R2', 'Test MSE', 'Test R2']

rf_results

```

Out[52]:

	Method	Training MSE	Training R2	Test MSE	Test R2
0	Random Forest	3133501584.973712	0.492041	3474153679.222308	0.519543

## Model Comparison: Linear regression vs. Random forest

In [53]:

```

# Creating rf_results data as a list
rf_results = ['Random Forest', 3133501584.973712, 0.492041, 3474153679.222308, 0.519543]

# Adding the rf_results to the lr_results DataFrame
lr_results.loc[len(lr_results)] = rf_results
df_models = lr_results
df_models

```

Out[53]:

	Method	Training MSE	Training R2	Test MSE	Test R2
0	Linear Regression	3215818093.812236	0.478697	3112173832.758111	0.569603
1	Random Forest	3133501584.973712	0.492041	3474153679.222308	0.519543

**Observation:** The Mean Squared Error(Test MSE) of the linear regression model is a bit lower than that of the Random Forest regression model. This means the amount of prediction errors made by the model is lower. Therefore, the linear regression model is slightly better at predicting house prices compared to the random forest model.